

```
In [35]: import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
from matplotlib import pyplot as plt
```

```
In [24]: df=pd.read_csv("https://raw.githubusercontent.com/dsrs scientist/dataset3/refs/heads/main/glass
df
```

Out[24]:

	0	1	2	3	4	5	6	7	8	9	10
0	1	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.00	0.0	1
1	2	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.00	0.0	1
2	3	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.00	0.0	1
3	4	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.00	0.0	1
4	5	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.00	0.0	1
...	...	...	...	...	...	...	...	...	...	...	...
209	210	1.51623	14.14	0.00	2.88	72.61	0.08	9.18	1.06	0.0	7
210	211	1.51685	14.92	0.00	1.99	73.06	0.00	8.40	1.59	0.0	7
211	212	1.52065	14.36	0.00	2.02	73.42	0.00	8.44	1.64	0.0	7
212	213	1.51651	14.38	0.00	1.94	73.61	0.00	8.48	1.57	0.0	7
213	214	1.51711	14.23	0.00	2.08	73.36	0.00	8.62	1.67	0.0	7

214 rows × 11 columns

```
In [25]: df.shape
```

Out[25]: (214, 11)

```
In [26]: df.isnull().sum()
```

```
Out[26]: 0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
dtype: int64
```

```
In [27]: df[10].value_counts()
```

```
Out[27]: 10
2      76
1      70
7      29
3      17
5      13
6       9
Name: count, dtype: int64
```

In [28]: `df.describe()`

Out[28]:

	0	1	2	3	4	5	6	7	8
<b>count</b>	214.000000	214.000000	214.000000	214.000000	214.000000	214.000000	214.000000	214.000000	214.000000
<b>mean</b>	107.500000	1.518365	13.407850	2.684533	1.444907	72.650935	0.497056	8.956963	0.175047
<b>std</b>	61.920648	0.003037	0.816604	1.442408	0.499270	0.774546	0.652192	1.423153	0.497219
<b>min</b>	1.000000	1.511150	10.730000	0.000000	0.290000	69.810000	0.000000	5.430000	0.000000
<b>25%</b>	54.250000	1.516522	12.907500	2.115000	1.190000	72.280000	0.122500	8.240000	0.000000
<b>50%</b>	107.500000	1.517680	13.300000	3.480000	1.360000	72.790000	0.555000	8.600000	0.000000
<b>75%</b>	160.750000	1.519157	13.825000	3.600000	1.630000	73.087500	0.610000	9.172500	0.000000
<b>max</b>	214.000000	1.533930	17.380000	4.490000	3.500000	75.410000	6.210000	16.190000	3.150000

Above statistics shows that data is across all attributes is not in same range, so will normalize the data first.

## Preparing Dataset

Adding meaningful column/attribute names

In [29]: `names=['Id', 'RI', 'Na', 'Mg', 'Al', 'Si', 'K', 'Ca', 'Ba', 'Fe', 'glass_type']`  
`df.columns=names`  
`df.head()`

Out[29]:

	Id	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	glass_type
<b>0</b>	1	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
<b>1</b>	2	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.0	0.0	1
<b>2</b>	3	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
<b>3</b>	4	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1
<b>4</b>	5	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.0	0.0	1

Removing unnecessary columns

In [22]: `df=df.drop("Id")`

```

-----
KeyError                                Traceback (most recent call last)
Cell In[22], line 1
----> 1 df=df.drop("Id")

File ~\anaconda3\Lib\site-packages\pandas\core\frame.py:5258, in DataFrame.drop(self, labels, axis, index, columns, level, inplace, errors)
    5110 def drop(
    5111     self,
    5112     labels: IndexLabel = None,
    (...)
    5119     errors: IgnoreRaise = "raise",
    5120 ) -> DataFrame | None:
    5121     """
    5122     Drop specified labels from rows or columns.
    5123     (...)
    5256     weight 1.0      0.8
    5257     """
    5258     ...

```

```
In [21]: df.head(3)
```

```
Out[21]:
```

	Id	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	glass_type
0	1	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
2	3	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
3	4	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1

## Checking outliers through Z-score

```
In [31]: z=abs(stats.zscore(df))  
  
#np.where(z>3)  
  
df=df[(z<3).all(axis=1)]  
  
#df.shape
```

## Separating Features and Label

```
In [32]: features=['RI', 'Na', 'Mg', 'Al', 'Si', 'K', 'Ca', 'Ba', 'Fe']  
label=['glass_type']  
  
X=df[features]  
  
Y=df[label]
```

```
In [33]: X.shape
```

```
Out[33]: (194, 9)
```

```
In [34]: type(X)
```


```
Out[34]: pandas.core.frame.DataFrame
```

## Data Visualization

```
In [36]: x2=X.values

for i in range(1,9):
    sns.distplot(x2[i])
    plt.xlabel(features[i])
    plt.show()
```

```
sns.distplot(x2[i])
```



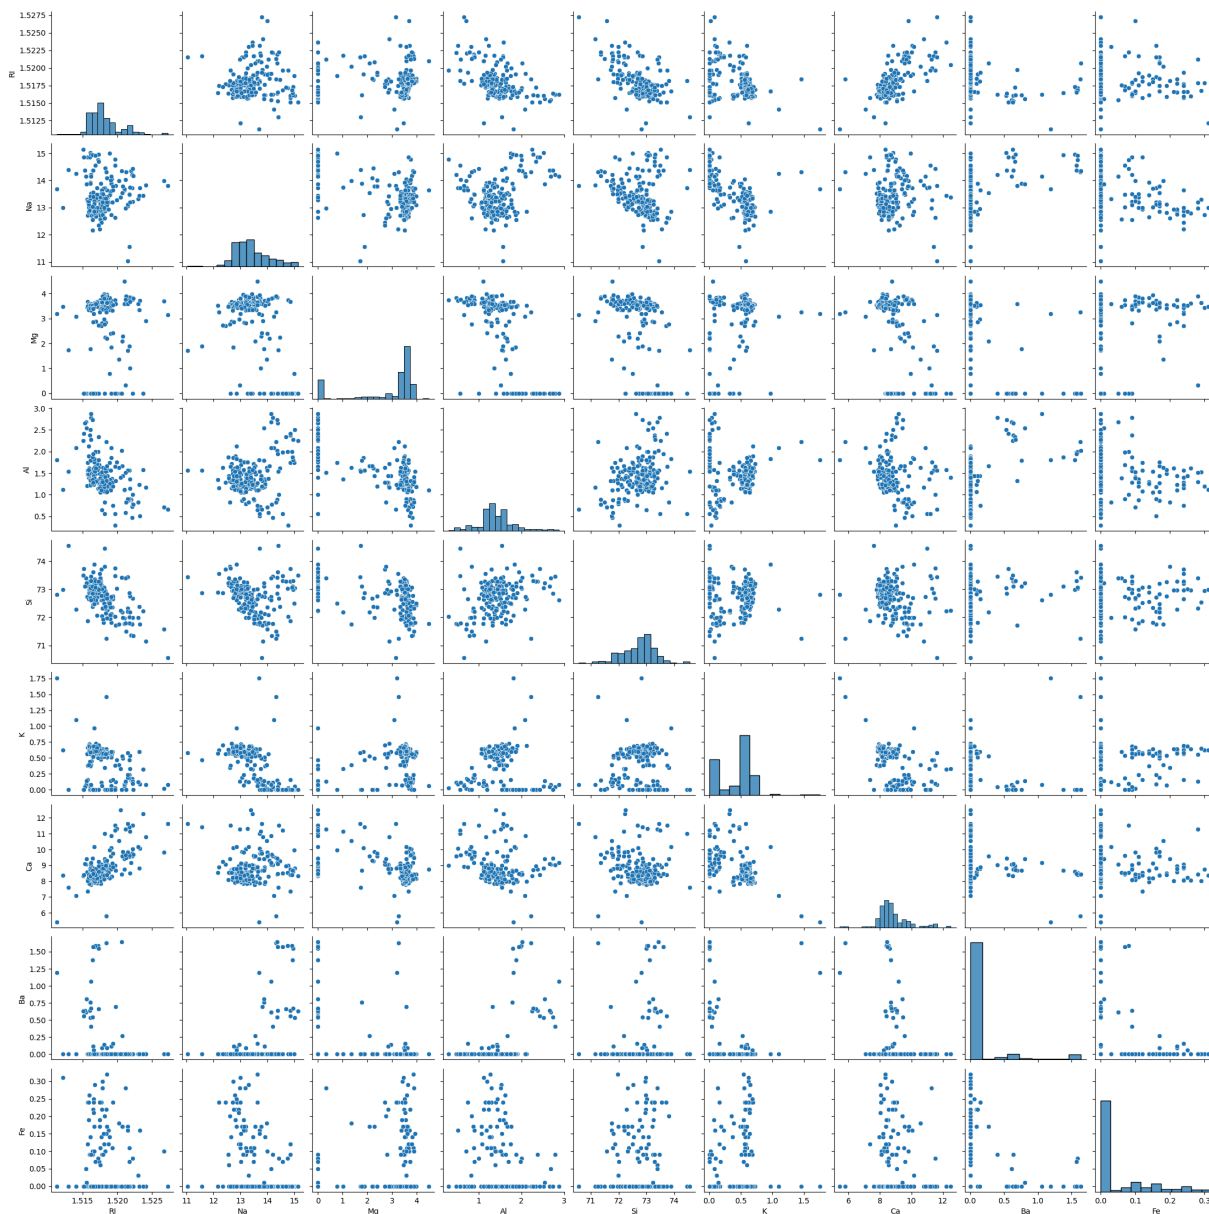
Above diagrams shows that our dataset is skewed either on positive side or negative side and data is not normalized.

```
In [38]: x2=pd.DataFrame(X)
plt.figure(figsize=(8,8))
sns.pairplot(data=x2)
plt.show()
```


C:\Users\Rashmi\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

self.\_figure.tight\_layout(\*args, \*\*kwargs)

<Figure size 800x800 with 0 Axes>



```
In [40]: correlation=X.corr()
figure(figsize=(15,15))
heatmap(correlation,cbar=True,square=True,annot=True,fmt='.1f',annot_kws={'size':15},xticklabels=
how())
```



Our Diagram shows correlation between different features conclusion:

1. RI and Ca have strong correlation between each other
2. Al and ba have intermediate correlation between each other

## Scaling tha data(1-0 range)

```
In [41]: # normalizing/Scaling the data

from sklearn.preprocessing import MinMaxScaler
scaler= MinMaxScaler()
#scaler.fit(X)
#X=scaler.transform(X)
#X=pd.DataFrame(X)
```

```
In [42]: X.head(2)
```

```
Out[42]:
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0
1	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.0	0.0

```
In [43]: Y.head(2)
```

```
Out[43]:
```

	glass_type
0	1
1	1

## Scaling the features

```
In [44]: from sklearn import preprocessing
X=preprocessing.scale(X)
```

## Visualizing Data after Preprocessing


```
In [46]: x2=X

from matplotlib import pyplot as plt
import seaborn as sns
for i in range(1,9):
    sns.distplot(x2[i])
    plt.xlabel(features[i])
    plt.show()
```

similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
sns.distplot(x2[i])
```



Above diagrams show that after preprocessing skewness is reduced and data is more normalized.

## Train Test Split

```
In [47]: from sklearn.model_selection import train_test_split
```

```
In [51]: X_train, X_test, y_train, y_test = train_test_split(X,Y,test_size=0.25, random_state=0,strati
```

```
In [52]: ##Flattening the array
```

```
y_train=y_train.values.ravel()  
y_test=y_test.values.ravel()
```

```
In [53]: print('Shape of X_train= ' + str(X_train.shape))  
print('Shape of X_test= ' + str(X_test.shape))  
print('Shape of y_train= ' + str(y_train.shape))  
print('Shape of y_test= ' + str(y_test.shape))
```

```
Shape of X_train= (145, 9)  
Shape of X_test= (49, 9)  
Shape of y_train= (145,)  
Shape of y_test= (49,)
```

## Applying Different Machine Models

### 1. KNN

```
In [54]: from sklearn.metrics import accuracy_score  
from sklearn.neighbors import KNeighborsClassifier
```

```
In [57]: Scores=[]  
  
for i in range(2,11):  
    knn=KNeighborsClassifier(n_neighbors=i)  
    knn.fit(X_train, y_train)  
    score=knn.score(X_test,y_test)  
    Scores.append(score)  
  
print(knn.score(X_train,y_train))  
print(Scores)
```

```
0.6896551724137931  
[0.7142857142857143, 0.6530612244897959, 0.7346938775510204, 0.7142857142857143, 0.673469387  
755102, 0.6530612244897959, 0.6938775510204082, 0.6938775510204082, 0.6938775510204082]
```

### 2. Decision Tree



In [58]: `from sklearn.tree import DecisionTreeClassifier`

```
Scores=[]

for i in range(1):
    tree=DecisionTreeClassifier(random_state=0)
    tree.fit(X_train,y_train)
    score=tree.score(X_test,y_test)
    Scores.append(score)

print(tree.score(X_train,y_train))
print(Scores)
```

```
1.0
[0.5510204081632653]
```

### 3. Logistic Regression

In [59]: `from sklearn.linear_model import LogisticRegression`

```
Scores=[]

for i in range(1):
    logistic =LogisticRegression(random_state=0, solver='lbfgs', multi_class='multinomial',max_iter=1000)
    logistic.fit(X_train, y_train)
    score = logistic.score(X_test,y_test)
    Scores.append(score)

print(logistic.score(X_train, y_train))
print(Scores)
```

```
0.7517241379310344
[0.6938775510204082]
```

### 4. SVC Classifier (Non-Linear Kernel)

In [61]: `from sklearn.svm import SVC`

```
Scores=[]

for i in range(1):
    svc=SVC(gamma='auto')
    svc.fit(X_train,y_train)
    score=svc.score(X_test, y_test)
    Scores.append(score)

print(svc.score(X_train, y_train))
print(Scores)
```

```
0.7517241379310344
[0.7551020408163265]
```

### 5. SVC Classifier (Linear Kernel)

```
In [62]: from sklearn.svm import LinearSVC

Scores=[]

for i in range(1):
    svc=LinearSVC(random_state=0)
    svc.fit(X_train, y_train)
    score=svc.score(X_test, y_test)
    Scores.append(score)

print(svc.score(X_train, y_train))
print(Scores)
```

```
0.7517241379310344
[0.6938775510204082]
```

C:\Users\Rashmi\anaconda3\Lib\site-packages\sklearn\svm\\_classes.py:32: FutureWarning: The default value of `dual` will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly to suppress the warning.

C:\Users\Rashmi\anaconda3\Lib\site-packages\sklearn\svm\\_base.py:1242: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

## 6. Random Forest

```
In [63]: from sklearn.ensemble import RandomForestClassifier

Scores=[]
Range=[10,20,30,50,70,80,100,120]

for i in range(1):
    forest=RandomForestClassifier(criterion='gini',n_estimators=10, min_samples_leaf=1, min_s
    forest.fit(X_train, y_train)
    score=forest.score(X_test, y_test)

print(forest.score(X_train, y_train))
print(score)
```

```
0.9724137931034482
0.7755102040816326
```

In [ ]:

## 7. Gradient Decent Tree Boosting

```
In [64]: from sklearn.ensemble import GradientBoostingClassifier

gd=GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1, random_state=
gd.fit(X_train, y_train)
score=gd.score(X_test, y_test)

print(gd.score(X_train, y_train))
print(score)
```

```
0.9724137931034482
0.6326530612244898
```

## Summary

Out of all above models:

1. Random forest is giving best result with:

```
Training accuracy: 0.9724137931034482  
Testing accuracy: 0.7755102040816326
```

But since it is overfitting we will choose next best model that is:

```
2. SVM (Non Linear Kernel)  
Training accuracy: 0.7517241379310344  
Testing accuracy: 0.7551020408163265
```

In [ ]: